


# Comp 428 – Fall 2018

Tutorial 1

Introduction to MPI




# Journey of Computing Tech

# Serial Computing

- ▶ Traditionally, software has been written for **serial** computation:
  - ▶ To be run on a single computer having a single Central Processing Unit (CPU)
  - ▶ A problem is broken into a discrete series of instructions
  - ▶ Instructions are executed one after another
  - ▶ Only one instruction may be executed at any moment in time

# Serial Computing




# Challenges

## ► Demand for edge computing



Exponential View



# Motivation for Parallel Computing

- ▶ Technology push

- ▶ Latest development happening on Hardware side


- ▶ Application pull

- ▶ Complex problems require computation on large-scale data
  - ▶ Sufficient performance available only through massive parallelism

# Parallel Computing

- ▶ In the simplest sense, **parallel computing** is the simultaneous use of multiple computing resources to solve a computational problem:
  - ▶ To be run using multiple CPUs
  - ▶ A problem is broken down to a series of instructions
  - ▶ Instructions from each part execute simultaneously on **different CPUs**


# Parallel Computing






# Question?

# Concurrency vs Parallelism




Concurrent: 2 queues, 1 vending machine



Parallel: 2 queues, 2 vending machines

# How to drive these horses?



# What needs to be done

- ▶ Initialization
- ▶ Work allocation
- ▶ Computation & Communication
- ▶ Termination



# Need of standards



<http://www.lordofthejars.com/2017/10/adding-chaos-on-openshift-cluster.html>

# Message passing Interface (MPI)


- ▶ MPI is a specification for the developers and users of message passing libraries.
- ▶ It is NOT a library - but rather the specification of what such a library should be.
- ▶ The goal of the Message Passing Interface is to provide a widely used standard for writing message passing programs. The interface attempts to be
  - ▶ Practical
  - ▶ Portable
  - ▶ Efficient
  - ▶ Flexible
- ▶ Interface specifications have been defined for C/C++ and Fortran programs.
- ▶ Complete MPI specification consists of 129 calls.

# Goals of the MPI standard

- ▶ MPI's prime goals are:
  - ▶ To provide source-code portability
  - ▶ To allow efficient implementations
- ▶ MPI also offers:
  - ▶ A great deal of functionality
  - ▶ Support for heterogeneous parallel architectures
  - ▶ Allow efficient and reliable Communication

# MPI Programming Model


- ▶ Originally, MPI was designed for distributed memory architectures, which were becoming increasingly popular at that time.





# MPI Programming Model (Cont)

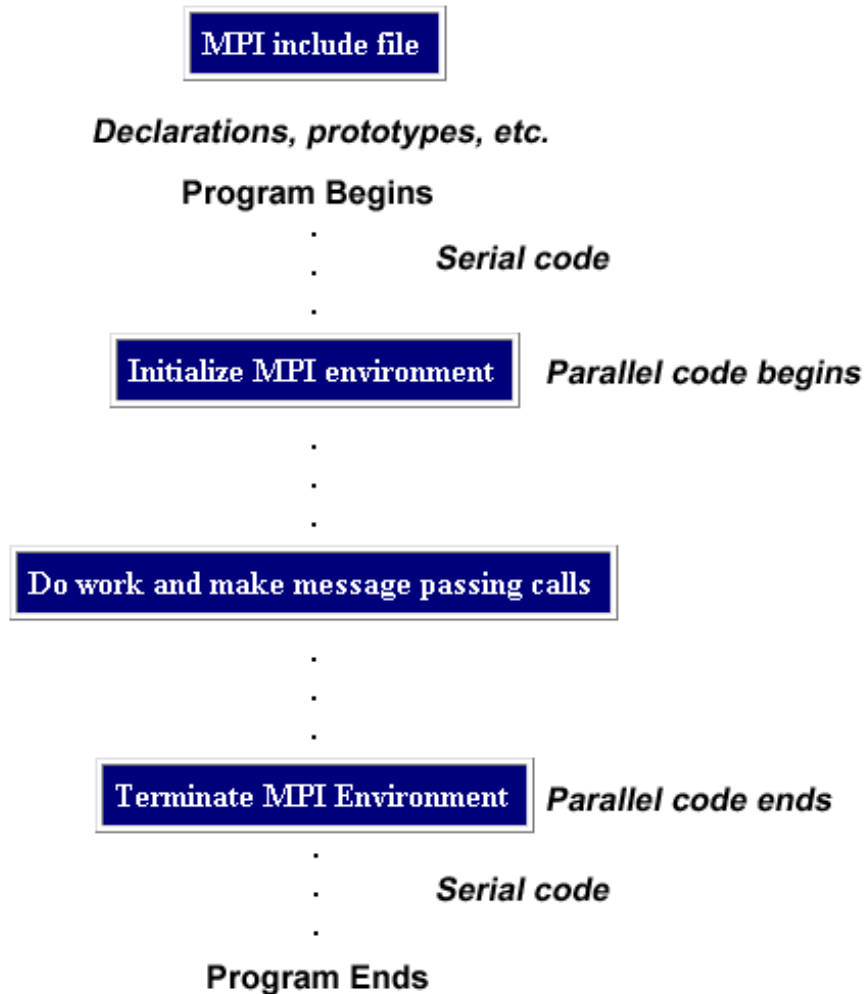
- ▶ As architecture trends changed, shared memory Symmetric Multiprocessors (SMPs) were combined over networks creating hybrid distributed memory / shared memory systems. MPI implementers adapted their libraries to handle both types of underlying memory architectures seamlessly.



# MPI Programming Model (Cont)

- ▶ Today, MPI runs on virtually any hardware platform:
  - ▶ Distributed Memory
  - ▶ Shared Memory
  - ▶ Hybrid
- ▶ All parallelism is explicit: the programmer is responsible for correctly identifying parallelism and implementing parallel algorithms using MPI constructs.

# General MPI Program Structure



# MPI is Simple

- ▶ Many parallel programs can be written using just these six functions
  - ▶ MPI\_Init
  - ▶ MPI\_Comm\_size
  - ▶ MPI\_Comm\_rank
  - ▶ MPI\_Send
  - ▶ MPI\_Recv
  - ▶ MPI\_Finalize

# MPI basic functions (subroutines)

- ▶ `MPI_INIT`: initialize MPI(MPI run time environment)
- ▶ `MPI_COMM_SIZE`: returns the number of MPI processes (important for decomposition)
- ▶ `MPI_COMM_RANK`: returns unique id for each processor (identify the PE) within the specified communicator
- ▶ `MPI_SEND` and `MPI_RECV`: MPI data is not shared but can be communicated. Each process has its own data. To communicate one process may send some data to another.
- ▶ `MPI_FINALIZE`: shuts down run time environment
- ▶ Note: This is crucial, failure to properly shutdown code could (will) cause other nodes in program to hang

# Serial Program

```
#include <stdio.h>

int main(){
    printf ("Welcome to COMP428
            Tutorial!\n");
    return 0;
}
```

# Serial to Parallel

```
#include <stdio.h>
#include "mpi.h"
int main(){
    MPI_Init (NULL, NULL);  /*start MPI*/
    printf ("Welcome to COMP428
            Tutorial!\n");
    MPI_Finalize();        /*shut down MPI*/
    return 0;
}
```

# MPI Build Scripts:

- Available scripts are listed below:

Language	Script Name	Underlying Compiler
C	<b>mpicc</b>	<b>gcc</b>
	<b>mpigcc</b>	<b>gcc</b>
	<b>mpiicc</b>	<b>Icc</b>
C++	<b>mpipgcc</b>	<b>pgcc</b>
	<b>mpiCC</b>	<b>g++</b>
	<b>mpig++</b>	<b>g++</b>
	<b>mpiicpc</b>	<b>icpc</b>
	<b>mpipgCC</b>	<b>pgCC</b>



# Format of MPI Calls:

- ▶ Programs must not declare variables or functions with names beginning with the prefix MPI\_

C binding	
Format	<code>rc = MPI_Xxxxxx(parameter,...)</code>
Example	<code>rc = MPI_Bsend(&amp;buf, count)</code>
Error Code	Returned as “rc”. MPI_SUCCESS if successful

# Communicators and Groups

- ▶ MPI uses objects called communicators and groups to define which collection of processes may communicate with each other.
- ▶ Most MPI routines require you to specify a communicator as an argument.
- ▶ Communicators and groups will be covered in more detail later. For now, simply use `MPI_COMM_WORLD` whenever a communicator is required - it is the predefined communicator that includes all of your MPI processes.

# Rank

- ▶ Within a communicator, every process has its own unique, integer identifier assigned by the system when the process initializes. A rank is sometimes also called a "task ID".
- ▶ Ranks are contiguous and begin at zero.
- ▶ Used by the programmer to specify the source and destination of messages. Often used conditionally by the application to control program execution (if rank=0 do this / if rank=1 do that).

# Error Handling:

- ▶ Most MPI routines include a return/error code parameter, as described in previous slide "Format of MPI Calls".
- ▶ However, according to the MPI standard, the default behavior of an MPI call is to abort if there is an error. This means you will probably not be able to capture a return/error code other than `MPI_SUCCESS` (zero).
- ▶ The standard does provide a means to override this default error handle. You can consult the error handling section of the MPI Standard located at <http://www.mpi-forum.org/docs/mpi-1.1-html/node148.html>.
- ▶ The types of errors displayed to the user are implementation dependent

# Basic Features of MPI Programs

- ▶ Calls may be roughly divided into four classes:
  - ▶ Calls used to initialize, manage, and terminate communications
  - ▶ Calls used to communicate between pairs of processors. (Pair communication)
  - ▶ Calls used to communicate among groups of processors. (Collective communication)
  - ▶ Calls to create data types.

# How to build a simple MPI program

- ▶ Have to include “mpi.h” .
- ▶ All MPI program must call MPI\_INIT as the first MPI call to initialize themselves.
- ▶ Call MPI\_COMM\_SIZE to get the number of processes that are running.
- ▶ Call MPI\_COMM\_RANK to determine there number between 0 and (size-1).
- ▶ Conditional process and general message passing can take place.
- ▶ Must call MPI\_FINALIZE as the last call to an MPI library routine.

**Thank You**

**Questions ?**